

Machine Learning in Public Health

Lecture 6: Linear Model Selection and Regularization (Lab)

Dr. Yang Feng

Best Subset Selection

```
library(ISLR)
library(tidyverse)
data("Credit")
credit_cate <- Credit %>% dplyr::select((-ID))
credit <- model.matrix(~., data = credit_cate) %>%
  as_tibble() %>%
  select(-(Intercept))
set.seed(0)
tr_ind <- sample(1:nrow(credit), 200)
credit_tr <- credit[tr_ind,]
credit_te <- credit[-tr_ind,]
```

```
library(leaps)
best_subset <- regsubsets(Balance ~., data = credit_tr)
##default only considers models size up to 8
best_subset <- regsubsets(Balance ~., data = credit_tr, nvmax = 11)
summary(best_subset)
```

```
## Subset selection object
## Call: regsubsets.formula(Balance ~ ., data = credit_tr, nvmax = 11)
## 11 Variables (and intercept)
##              Forced in Forced out
## Income                FALSE      FALSE
## Limit                  FALSE      FALSE
## Rating                 FALSE      FALSE
## Cards                  FALSE      FALSE
## Age                   FALSE      FALSE
## Education              FALSE      FALSE
## GenderFemale           FALSE      FALSE
## StudentYes             FALSE      FALSE
## MarriedYes            FALSE      FALSE
## EthnicityAsian        FALSE      FALSE
## EthnicityCaucasian    FALSE      FALSE
## 1 subsets of each size up to 11
## Selection Algorithm: exhaustive
##           Income Limit Rating Cards Age Education GenderFemale StudentYes
## 1  ( 1 )  " "    " "    "*"    " "    " "    " "    " "
## 2  ( 1 )  "*"    " "    "*"    " "    " "    " "    " "
```

```
## 3 ( 1 ) "*" "*" " " " " " " " " " " "*"
## 4 ( 1 ) "*" "*" " " "*" " " " " " " "*"
## 5 ( 1 ) "*" "*" " " "*" "*" " " " " "*"
## 6 ( 1 ) "*" "*" "*" "*" "*" " " " " "*"
## 7 ( 1 ) "*" "*" "*" "*" "*" " " " " "*"
## 8 ( 1 ) "*" "*" "*" "*" "*" " " " " "*"
## 9 ( 1 ) "*" "*" "*" "*" "*" " " " " "*"
## 10 ( 1 ) "*" "*" "*" "*" "*" " " " " "*"
## 11 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*"
##      MarriedYes EthnicityAsian EthnicityCaucasian
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " " " " "
## 8 ( 1 ) "*" " " " "
## 9 ( 1 ) "*" " " "*"
## 10 ( 1 ) "*" "*" "*"
## 11 ( 1 ) "*" "*" "*"

```

```
##now considers all 11 variables
best_subset_sum <- summary(best_subset)

best_subset_sum$rsq

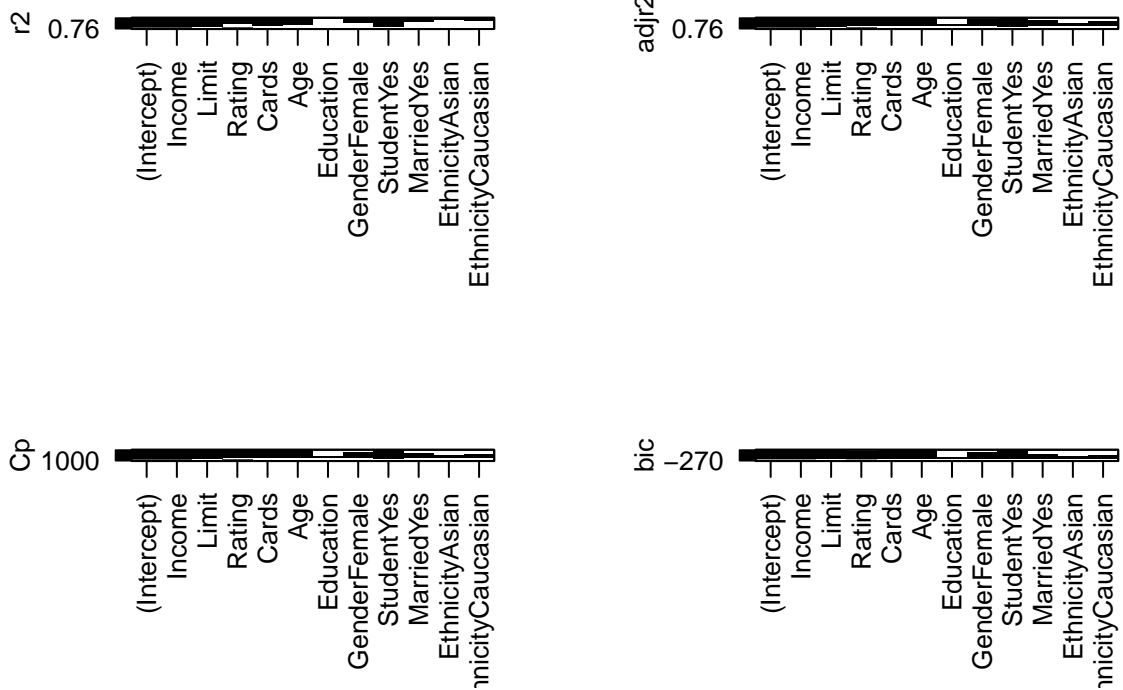
```

```
## [1] 0.7590176 0.8813666 0.9574599 0.9616792 0.9625211 0.9625861 0.9626233
## [8] 0.9626522 0.9626658 0.9626853 0.9626897

```


```
par(mfrow = c(2,2))
plot(best_subset, scale = "r2")
plot(best_subset, scale = "adjr2")
plot(best_subset, scale = "Cp")
plot(best_subset, scale = "bic")

```



Optimal model for each size: four methods

```
measures <- c("rsq", "adjr2", "cp", "bic")
our_names <- c("R2", "Adjusted R2", "Cp", "BIC")
size_seq <- 1:length(best_subset_sum$rsq)
my_plots <- NULL
for(mea_ind in seq_along(measures)){
  dat <- data.frame(d = size_seq, val = best_subset_sum[[measures[mea_ind]]])
  my_plots[[mea_ind]] <- ggplot(dat, mapping = aes(x = d, y = val)) + geom_point() + geom_line() +
    ggtitle(our_names[mea_ind])
}
library(egg)
grid.arrange(grobs = my_plots, ncol = 2)
```



Coefficients corresponding to optimal models

```
coef(best_subset, 1:11)
```

```
## [[1]]
## (Intercept)      Rating
## -354.238144      2.501169
##
## [[2]]
## (Intercept)      Income      Rating
## -548.260759     -7.684998      4.002145
##
## [[3]]
## (Intercept)      Income      Limit      StudentYes
## -456.4610300     -7.8639273      0.2723742      429.2673608
##
## [[4]]
## (Intercept)      Income      Limit      Cards      StudentYes
## -523.0518688     -7.7925423      0.2712276      23.0613375      437.2257637
##
## [[5]]
## (Intercept)      Income      Limit      Cards      Age      StudentYes
## -474.4151238     -7.6386077      0.2695746      23.6202781     -0.8761518      435.8910167
##
```

```
## [[6]]
## (Intercept)      Income      Limit      Rating      Cards      Age
## -483.1236387    -7.6550380    0.2425861    0.4057927    21.4430273    -0.8818302
## StudentYes
## 434.7457520
##
## [[7]]
## (Intercept)      Income      Limit      Rating      Cards      Age
## -479.7062751    -7.6587166    0.2432506    0.3961782    21.4556382    -0.8844975
## GenderFemale StudentYes
## -5.9442936    434.5896004
##
## [[8]]
## (Intercept)      Income      Limit      Rating      Cards      Age
## -476.0559413    -7.6549983    0.2430891    0.3995257    21.4452607    -0.8982854
## GenderFemale StudentYes MarriedYes
## -5.6678976    433.4931082    -5.4990089
##
## [[9]]
## (Intercept)      Income      Limit      Rating
## -478.2999819    -7.6539473    0.2430925    0.3992281
## Cards      Age      GenderFemale StudentYes
## 21.4787522    -0.8909104    -5.6975997    433.3446404
## MarriedYes EthnicityCaucasian
## -5.5205462    3.6022736
##
## [[10]]
## (Intercept)      Income      Limit      Rating
## -482.4570451    -7.6484920    0.2412729    0.4253337
## Cards      Age      GenderFemale StudentYes
## 21.4309499    -0.8814066    -5.7321331    433.2456953
## MarriedYes EthnicityAsian EthnicityCaucasian
## -6.0516622    6.2314770    6.8553322
##
## [[11]]
## (Intercept)      Income      Limit      Rating
## -477.5268904    -7.6464664    0.2415817    0.4201142
## Cards      Age      Education      GenderFemale
## 21.4194258    -0.8833370    -0.3266788    -5.7043018
## StudentYes MarriedYes EthnicityAsian EthnicityCaucasian
## 433.3529352    -6.1769790    6.3441993    6.8727948
```

```
coef(best_subset, 4)
```

```
## (Intercept)      Income      Limit      Cards      StudentYes
## -523.0518688    -7.7925423    0.2712276    23.0613375    437.2257637
```

```
best_ind <- which.min(best_subset_sum$bic)
best_coef <- coef(best_subset, best_ind)
tr_x <- credit_tr %>% select(names(best_coef)[-1])
tr_pred <- cbind(1, as.matrix(tr_x)) %*% best_coef
tr_error_best <- mean((tr_pred - credit_tr$Balance)^2)
te_x <- credit_te %>% select(names(best_coef)[-1])
```

```
te_pred <- cbind(1, as.matrix(te_x)) %*% best_coef
te_error_best <- mean((te_pred - credit_te$Balance)^2)
tr_error_best
```

```
## [1] 9033.365
```

```
te_error_best
```

```
## [1] 10837.48
```

Forward Stepwise Selection

```
forward_fit <- regsubsets(Balance ~., data = credit_tr, method = "forward", nvmax = 11)
forward_sum <- summary(forward_fit)
best_ind <- which.min(forward_sum$bic)
best_coef <- coef(forward_fit, best_ind)
```

```
tr_x <- credit_tr %>% select(names(best_coef)[-1])
tr_pred <- cbind(1, as.matrix(tr_x)) %*% best_coef
tr_error_forward <- mean((tr_pred - credit_tr$Balance)^2)
te_x <- credit_te %>% select(names(best_coef)[-1])
te_pred <- cbind(1, as.matrix(te_x)) %*% best_coef
te_error_forward <- mean((te_pred - credit_te$Balance)^2)
tr_error_forward
```

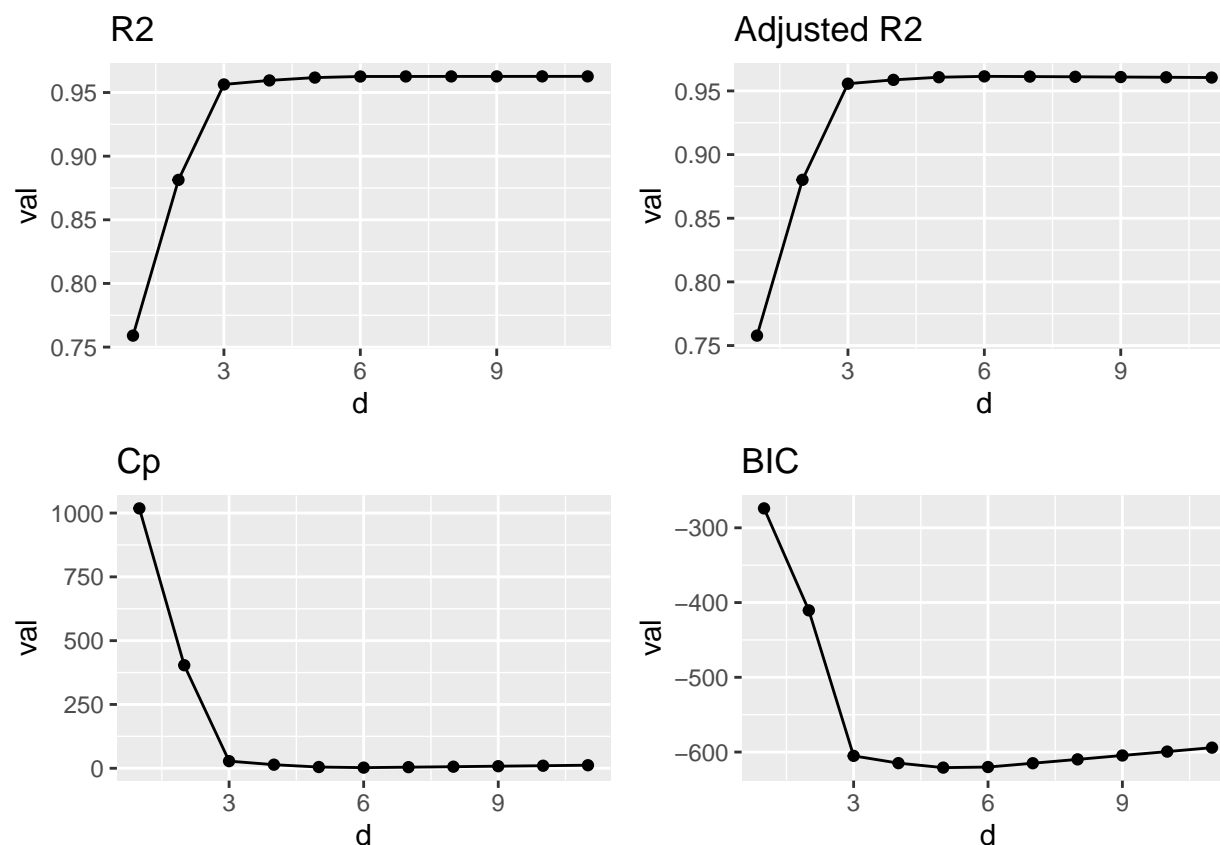
```
## [1] 9020.523
```

```
te_error_forward
```

```
## [1] 10702.69
```

Optimal model for each size: four methods

```
size_seq <- 1:length(forward_sum$rsq)
my_plots <- NULL
for(mea_ind in seq_along(measures)){
  dat <- data.frame(d = size_seq, val = forward_sum[[measures[mea_ind]]])
  my_plots[[mea_ind]] <- ggplot(dat, mapping = aes(x = d, y = val)) + geom_point() + geom_line() +
    ggtitle(our_names[mea_ind])
}
grid.arrange(grobs = my_plots, ncol = 2)
```



Backward Stepwise Selection

```
backward_fit <- regsubsets(Balance ~., data = credit_tr, method = "backward", nvmax = 11)
backward_sum <- summary(backward_fit)
best_ind <- which.min(backward_sum$bic)
best_coef <- coef(backward_fit, best_ind)
```

```
tr_x <- credit_tr %>% select(names(best_coef)[-1])
tr_pred <- cbind(1, as.matrix(tr_x)) %*% best_coef
tr_error_backward <- mean((tr_pred - credit_tr$Balance)^2)
te_x <- credit_te %>% select(names(best_coef)[-1])
te_pred <- cbind(1, as.matrix(te_x)) %*% best_coef
te_error_backward <- mean((te_pred - credit_te$Balance)^2)
tr_error_backward
```

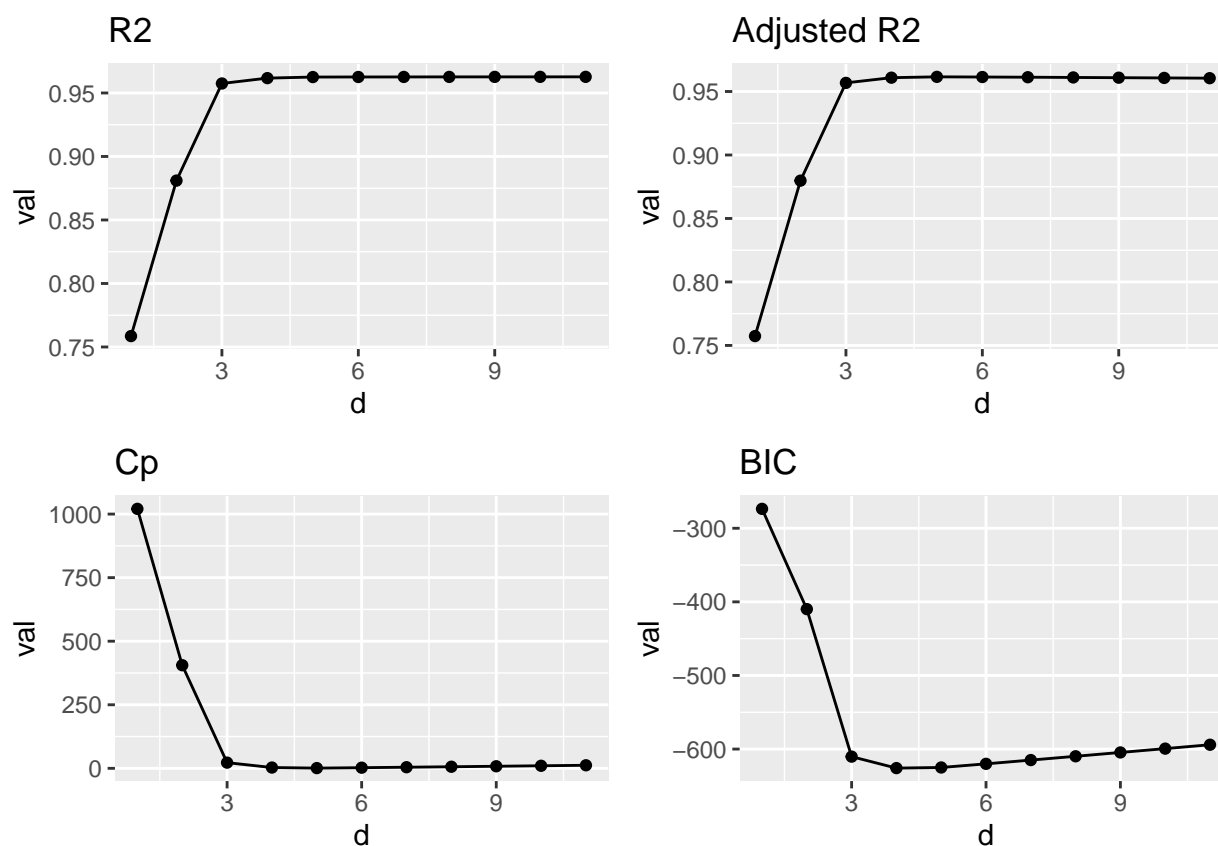
```
## [1] 9033.365
```

```
te_error_backward
```

```
## [1] 10837.48
```

Optimal model for each size: four methods

```
size_seq <- 1:length(backward_sum$rsq)
my_plots <- NULL
for(meas_ind in seq_along(measures)){
  dat <- data.frame(d = size_seq, val = backward_sum[[measures[meas_ind]]])
  my_plots[[meas_ind]] <- ggplot(dat, mapping = aes(x = d, y = val)) + geom_point() + geom_line() +
    ggtitle(our_names[meas_ind])
}
grid.arrange(grobs = my_plots, ncol = 2)
```



Ridge solution path for Credit data

```
library(glmnet)
library(caret)
library(plotmo)
x_tr <- as.matrix(credit_tr[, -12])
y_tr <- credit_tr[, 12, drop = T]
x_te <- as.matrix(credit_te[, -12])
y_te <- credit_te[, 12, drop = T]

std_fit <- preProcess(x_tr, method = c("center", "scale"))
```




```

x_tr_std <- predict(std_fit, x_tr)
x_te_std <- predict(std_fit, x_te)

fit_ridge <- glmnet(x_tr_std, y_tr, alpha = 0)

plot_glmnet(fit_ridge)

```



```

set.seed(0)
# with standardization
cv_fit_ridge <- cv.glmnet(x_tr, y_tr, alpha = 0)
tr_pred <- predict(cv_fit_ridge, newx = x_tr)
te_pred <- predict(cv_fit_ridge, newx = x_te)
tr_error_ridge <- mean((tr_pred - y_tr)^2)
te_error_ridge <- mean((te_pred - y_te)^2)
tr_error_ridge

```

```
## [1] 14812.97
```

```
te_error_ridge
```

```
## [1] 16207.82
```

```

# without standardization
cv_fit_ridge <- cv.glmnet(x_tr_std, y_tr, alpha = 0)
tr_pred <- predict(cv_fit_ridge, newx = x_tr_std)
te_pred <- predict(cv_fit_ridge, newx = x_te_std)
tr_error_ridge <- mean((tr_pred - y_tr)^2)
te_error_ridge <- mean((te_pred - y_te)^2)
tr_error_ridge

```


```
## [1] 14812.97
```

```
te_error_ridge
```

```
## [1] 16207.82
```

lasso Solution Path

```
fit_lasso <- glmnet(x_tr_std, y_tr, alpha = 1)
plot_glmnet(fit_lasso)
```



```
cv_fit_lasso <- cv.glmnet(x_tr, y_tr)
tr_pred <- predict(cv_fit_lasso, newx = x_tr)
te_pred <- predict(cv_fit_lasso, newx = x_te)
tr_error_lasso <- mean((tr_pred - y_tr)^2)
te_error_lasso <- mean((te_pred - y_te)^2)
tr_error_lasso
```

```
## [1] 10361.33
```

```
te_error_lasso
```

```
## [1] 11532.49
```

```
df <- data.frame(methods = c("Best Subset", "Forward", "Backward", "Ridge", "Lasso"),
                 test_error = c(te_error_best, te_error_forward, te_error_backward, te_error_ridge, te_
df
```

```
##      methods test_error
## 1 Best Subset  10837.48
## 2   Forward   10702.69
## 3  Backward   10837.48
## 4    Ridge   16207.82
## 5    Lasso   11532.49
```