

Machine Learning in Public Health

Lecture 8: Bagging, Random Forest, Boosting

Dr. Yang Feng

Advantages and disadvantages of trees

- Trees are very easy to explain to people. In fact, they are even easier to explain than linear regression!
- Why? No mathematical formula needed in the communication
- Some people believe that decision trees more closely mirror human decision-making than do the regression and classification approaches seen in previous chapters.
- Trees can be displayed graphically, and are easily interpreted even by a non-expert (especially if they are small).
- **Unfortunately**, trees generally do not have the great predictive accuracy.
- This disadvantage motivates the ensemble methods such as bagging, random forests and boosting.

Bagging

- The decision trees suffer from high variance.
- **Bootstrap aggregation**, or **bagging**, is a general-purpose procedure for reducing the variance of a statistical learning method; we introduce it here because it is particularly useful and frequently used in the context of decision trees.
- In this approach we generate B different *bootstrapped* training data sets. We then train our method on the b th bootstrapped training set in order to get $\hat{f}^{*b}(x)$ (not pruned), and average all the predictions:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x) .$$

Bagging

- Averaging these B trees reduces the variance. (Why?)
- B is not a critical parameter with bagging; a very large value of B will not lead to overfitting.
- In practice we use B sufficiently large that the error has settled down.
- For a given test observation in *classification*, we can record the class predicted by each of the B trees, and take a majority vote: the overall prediction is the most commonly occurring majority class among the B predictions.

Out-of-bag Error Estimation

- For a bagged model, we can estimate the test error without doing CV or a validation set approach
- Each bagged tree makes use of a part of the original observations
- The remaining observations not used to fit a given bagged tree are referred to as the **out-of-bag** (OOB) observations
- We can predict the response for the i -th observation using each of the trees in which that observation was OOB

Out-of-bag Error Estimation

- This will yield a little more than $B/3$ predictions for the i th observation
- To obtain a single prediction for the i -th observation, we can average these predicted responses (if regression is the goal) or can take a majority vote (if classification is the goal). This leads to a single OOB prediction for the i -th observation
- An OOB prediction can be obtained in this way for each of the n observations, from which the overall OOB MSE (for a regression problem) or classification error (for a classification problem) can be computed

Variable Importance Measures

- Although the collection of bagged trees is much more difficult to interpret than a single tree, one can obtain an overall summary of the importance of each predictor using
 - *the RSS (for bagging regression trees)*
 - *the Gini index (for bagging classification trees)*
- In the case of bagging regression trees, we can record the total amount that the RSS is decreased due to splits over a given predictor, averaged over all B trees. A large value indicates an important predictor
- For bagging classification trees, we can add up the total amount that the Gini index is decreased by splits over a given predictor, averaged over all B trees.

Variable Importance Measure

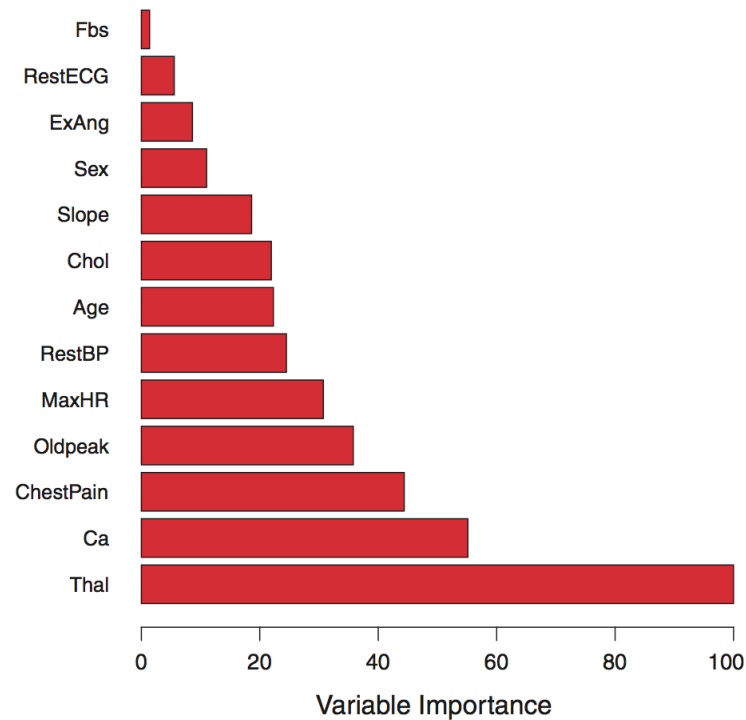


FIGURE 8.9. A variable importance plot for the **Heart** data. Variable importance is computed using the mean decrease in Gini index, and expressed relative to the maximum.


Random Forests

- Bagging constructs trees that are too “similar” (why?), so it probably does not reduce the variance as much as we wish to.
- **Random forests** provide an improvement over bagged trees by way of a small tweak that *decorrelates* the trees.
- As in bagging, we build a number of decision trees on bootstrapped training samples.
- But when building these decision trees, each time a split in a tree is considered, a random sample of m predictors is chosen as split candidates from the full set of p predictors. The split is allowed to use **only one** of those m predictors.

Random Forests

- Bagging is a special case of random forest when $m = p$
- If you don't want to spend extra efforts on m , $mtry = \sqrt{p}$ is a canonical choice for classification and $mtry = p/3$ is a good choice for regression.
- As with bagging, random forests will **not overfit** if we increase B , so in practice we use B sufficiently large for the error rate to have settled down. The default choice in the `randomForest` function is $ntree = 500$.
- Random Forest is a really good off-the-shelf algorithm!

Test and OOB Error



Some questions

- Q: In random forest algorithms, we restrict our attention to randomly selected m out of p features in each split. Now we change this procedure to restriction to the first m features (i.e., X_1, \dots, X_m) in every split. Do you expect the new procedure to work well? And why?
- Q: Is random forest always a better algorithm compared to decision trees?
- Q: What are the sources of randomness random forests have?
Hint: 2.
- Q: (brainstorm, no clear answer): Can you think of some ways to combine linear regression and decision trees / random forest?

Boosting

- Like bagging, **boosting** is a general approach that can be applied to many statistical learning methods for regression or classification.

Algorithm 8.2 *Boosting for Regression Trees*

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
2. For $b = 1, 2, \dots, B$, repeat:
 - (a) Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to the training data (X, r) .
 - (b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x). \quad (8.10)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i). \quad (8.11)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x). \quad (8.12)$$

In the boosting tree algorithm

- We use cross-validation to select B .
- The shrinkage parameter λ , a small positive number. This controls the rate at which boosting learns. Typical values are 0.01 or 0.001.
- The number d of splits in each tree, which controls the complexity of the boosted ensemble. Often $d = 1$ works well, in which case each tree is a *stump*, consisting of a single split. More generally d is the interaction depth.
- Boosting is a *slow* learner.